

MARCUS A. OKONKWO

Product Data Science · Experimentation & Causal Inference

Oakland, CA · (510) 555-0139 · marcus.okonkwo@gmail.com · linkedin.com/in/maokonkwo

SUMMARY

Eight years designing experiments and defining metrics for products used by hundreds of millions of people. My work sits where causal inference meets product judgment: I am usually the person arguing that the statistically significant result is not the interesting one. Since 2023 I have focused on measuring generative-AI features, where the honest answer is often that the metric does not exist yet and has to be invented.

PROFESSIONAL EXPERIENCE

Staff Product Data Scientist

Feb 2023 – Present

Figma – San Francisco, CA

- Founding data scientist on Figma AI. Built the measurement framework for generative features from nothing; defined *assisted edit acceptance*, *rework rate*, and a suggestion-quality index blended from implicit signals and a weekly 500-sample human rating panel.
- Ran the experiment program for the AI beta (23 tests over 14 months). Established a Bayesian readout for low-traffic surfaces after showing frequentist tests were underpowered for 80% of the beta cohort, and documented when each approach applies.
- Found that a headline +12% AI-suggestion acceptance rate coincided with a 4% drop in file-completion – traced it to suggestions that were easy to accept and hard to keep. The recommendation to hold the launch was adopted; the reworked version shipped two quarters later with both metrics positive.
- Own the weekly product metrics review for a 60-person org; wrote the metric charter that resolved a long-running WAU definition dispute between Growth and Product.

Data Scientist → Senior Data Scientist

Sep 2019 – Feb 2023

Meta – Menlo Park, CA

- Supported Marketplace integrity and later Reels ranking. Designed geo-level and cluster-randomized experiments where user-level assignment leaked across the social graph; the geo methodology note became internal reference material.
- Led the analysis behind a seller-verification rollout: combined diff-in-diff on staged geo launches with a survey instrument, sizing a 17% reduction in reported scam contacts against a 2.3% listing-volume cost.
- Built long-run holdouts for Reels ranking to separate durable engagement gains from novelty; roughly 40% of short-term wins did not survive 90 days, which changed how the team set launch bars.
- Mentored four junior data scientists and two PhD interns; two intern projects shipped to production.

Research Analyst

Jul 2018 – Aug 2019

Nielsen – New York, NY

- Built panel-weighting and audience-projection models for streaming clients; automated a monthly reporting pipeline in Python and SQL that replaced a 3-day manual process.

EDUCATION

M.A. Economics – Duke University · *fields: econometrics, industrial organization*

2018

B.S. Mathematics – Morehouse College · *minor in computer science*

2016

TECHNICAL SKILLS

LANGUAGES	SQL (Presto, Spark SQL, BigQuery), Python, R, some Stan
LIBRARIES	pandas, statsmodels, scikit-learn, PyMC, DoWhy, EconML, PyTorch (working knowledge)
METHODS	Randomized experiments (user, cluster, geo, switchback), Bayesian A/B testing, difference-in-differences, synthetic control, instrumental variables, regression discontinuity, power analysis and MDE planning
TOOLING	dbt, Airflow, Spark, Tableau, Mode, Streamlit, Git
QUALITATIVE	Survey design and weighting, structured user-interview coding (paired with UXR)

SELECTED TALKS & WRITING

- "When Your Treatment Leaks: Cluster Randomization in Social Products" – CODE@MIT, 2022.
- "Measuring Features That Do Not Have a Right Answer" – Figma engineering blog, 2024.