

Priya Ramanathan

PRODUCT DATA SCIENTIST

CONTACT

San Francisco, CA 94110
(415) 555-0182
priya.ramanathan@gmail.com
linkedin.com/in/priyaramanathan
github.com/pramanathan

TECHNICAL

Languages

SQL (Presto, BigQuery, Snowflake), Python, R, Scala (basic)

Python stack

pandas, NumPy, statsmodels, scikit-learn, PyMC, Polars

Data / BI

dbt, Airflow, Spark, Looker, Mode, Hex, Amplitude

Methods

A/B & switchback testing, CUPED, sequential testing, diff-in-diff, synthetic control, uplift modeling, survey design

EDUCATION

M.S. Statistics

Stanford University · 2017

B.A. Mathematics & Economics

UCLA · 2015 · summa cum laude

EXTRAS

Co-author, "Variance Reduction in Low-Traffic Marketplaces" — CODE@MIT 2023 (poster)
Mentor, Women in Data Science Bay Area
Conversational Tamil & Hindi

PROFILE

Product data scientist with 8 years building measurement from the ground up at consumer and developer-facing companies. I own the boring, load-bearing parts — metric definitions, experiment design, the dashboards people actually trust — and I have spent the last two years measuring LLM-powered features where the outcome variable is fuzzy and the sample is small. Comfortable being the first analyst on a surface with no instrumentation and no baseline.

EXPERIENCE

Senior Data Scientist, Product

Mar 2022 – Present

Stripe · San Francisco, CA

- Data science lead for Stripe's AI support assistant. Designed the offline/online eval split and the human-rated quality rubric that became the launch gate; caught a 6pt containment regression in staged rollout that offline scores had missed entirely.
- Rebuilt the Checkout conversion metric after finding the legacy definition double-counted retried sessions — restated 18 months of history, then wrote the migration doc that got 40+ dashboards moved over in one quarter.
- Designed and ran 60+ experiments on Checkout and Link. Introduced CUPED and sequential stopping to the team's shared framework, cutting median time-to-decision from 21 to 12 days.
- Built the merchant-activation funnel model used in quarterly planning; the "first-10-payments" metric it surfaced is now an org-level goal.
- Ran a paired qual study with UXR on failed onboarding sessions — the four themes it produced drove the 2024 onboarding rewrite, which lifted 30-day activation 9.4%.

Data Scientist II → III

Jul 2019 – Mar 2022

Airbnb · San Francisco, CA

- Owned analytics for Guest Search Ranking. Ran switchback tests to handle marketplace interference after demonstrating that naïve user-level splits overstated ranking wins by ~35%.
- Defined "quality booking," the guest-side counter-metric adopted across Search, Pricing and Trust; wrote the spec and the dbt models behind it.
- Built the self-serve experiment readout in Superset that removed roughly 15 analyst-hours/week of one-off requests.

Data Analyst

Aug 2017 – Jun 2019

Wayfair · Boston, MA

- Supported merchandising and pricing teams; built the first automated price-elasticity report for the furniture category.
- Migrated recurring SQL reporting into an Airflow + Looker stack, retiring 30+ spreadsheet trackers.